

# Machine Learning for Signal Processing

*[5LSL0]*

Rik Vullings  
Ruud van Sloun  
Nishith Chennakeshava  
Hans van Gorp

## **Assignment 2: Nonlinear Models**

April 2022

# Learning nonlinear functions for regression and classification

---

In this assignment we will explore some of the basic elements of machine learning in the context of optimal parameter estimation for nonlinear function approximation, regression and classification.

---

## Linear models

Consider the following regression network:

$$f(\mathbf{x}; \mathbf{w}, b) = \mathbf{w}^T \mathbf{x} + b, \quad (1.1)$$

which we will use to map a set of input vectors  $\mathbf{x}_i$  to a set of outputs  $y_i$ . The scalar outputs  $y_i$  are continuous variables stored in a vector  $y$ , and the input vectors  $\mathbf{x}_i$  are stored in a matrix  $X$  (see lecture slides).

Let us assume a Gaussian distribution of prediction errors with a unity covariance matrix, i.e.  $p(y_i; \theta)$  follows a normal distribution with mean  $f(\mathbf{x}_i; \theta)$  and variance 1.

**Q1: Derive the negative log-likelihood cost function  $J(\mathbf{x}, y; \theta)$  that will yield a maximum likelihood estimator of  $\mathbf{w}$  and  $b$ .**

**Q2: Given this cost function, derive an expression for the optimal weights  $\mathbf{w}$  and bias  $b$ . (Hint: Derive and equate the gradient of the cost function with respect to the parameters to zero.)**

We will now deploy our regression model and learn to predict outputs  $f(\mathbf{x}; \mathbf{w}, b)$  that match those produced by some unknown process  $f^*(\mathbf{x})$ .

Consider the inputs  $\mathbf{x} : \{[0, 0], [0.1, 1], [1, 0.2], [1, 1]\}$   
and process outputs  $f^*(\mathbf{x}) : \{0, 0.41, 0.18, 0.5\}$

**Q3: Derive the optimal parameters  $\mathbf{w}$  and  $b$  that describe this process. Is the process well-described by our regression model?**

For the same inputs, we now measure our process with some arbitrary noisy sensor, and obtain the following outputs:  $y : \{-0.0416, 0.3610, 0.1222, 0.4733\}$ .

**Q4: Recalculate the optimal parameters given these outputs. What happens? What would you do to obtain better estimates given such a noisy sensor?**

When the noise covariance structure of our measurements is not unity, our maximum-likelihood derived cost function changes.

**Q5: Derive the optimal negative log-likelihood cost criterion for a diagonal covariance matrix with  $\sigma_0, \dots, \sigma_i, \dots, \sigma_N$  on the diagonal. How does the balance between the variances of the parameter estimates play a role in the cost function?**

Now let us consider learning a new function, the XOR: given a 2-element input vector with binary entries, the output will be 1 if and only if one of the two elements in the input vector is 1, and the other 0.

**Q6: Calculate the optimal parameters of our regression model based on the four possible inputs and outputs of the XOR function.**

Notably, our model is unable to learn and perform the (rather simple) XOR function. In fact, the same holds for many other interesting functions that we would like *a machine* to learn. As we shall see, applying a nonlinear transformation into a new domain greatly increases the flexibility of our systems.

## Nonlinear functions

Before we proceed with nonlinear models, we define several nonlinear functions that are widely used in machine learning:

- **ReLU** (Rectified linear unit):  $f(x) = \max(0, x)$ .
- **Sigmoid**:  $f(x) = \sigma(x) = 1 / (1 + \exp(-x))$ .
- **Softmax**:  $f(\mathbf{x})_j = \frac{\exp(x_j)}{\sum_{i=0}^{K-1} \exp(x_i)}$  for  $j = 0, \dots, K - 1$ , where  $K$  is the size of vector  $\mathbf{x}$ .

**Q7: Derive compact expressions for the derivatives of the above nonlinearities with respect to their inputs.**

As for adaptive filters, training nonlinear regression models to fulfil a specific task is done through gradient-based learning. These derivatives will come in handy at that point.

**Q8: What happens to the gradients of the above three nonlinearities when the input values  $x \gg 0$ ?**

## Shallow (i.e. not deep...) nonlinear models

Consider a nonlinear model  $f(\mathbf{x}; \theta)$ :

$$f(\mathbf{x}; \theta) = f^{(2)}(\mathbf{h}; \mathbf{w}^{(2)}, b^{(2)}), \quad (1.2)$$

where

$$\mathbf{h} = f^{(1)}(\mathbf{x}; \mathbf{W}^{(1)}, \mathbf{b}^{(1)}) = \max(0, \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}), \quad (1.3)$$

nonlinearly maps  $\mathbf{x}$  into a new space  $\mathbf{h}$  through a ReLU.

The complete network is then:

$$f(\mathbf{x}; \mathbf{W}^{(1)}, \mathbf{b}^{(1)}, \mathbf{w}^{(2)}, b^{(2)}) = \left(\mathbf{w}^{(2)}\right)^T \max(0, \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) + b^{(2)}. \quad (1.4)$$

In the lectures, we saw that it is indeed possible to replicate an XOR function with this nonlinear model.

**Q9:** We sequentially apply two functions:  $f^{(1)}$  and  $f^{(2)}$ ; why should there be a nonlinearity in the first? Doesn't a linear mapping to a new space  $\mathbf{h}$  suffice?

**Q10:** Given the solution presented in the lecture slides (2-13), reproduce the plots on slide 2-14 in e.g. Python, showing how the input is mapped into the latent space  $\mathbf{h}$ . Also plot the decision boundary  $f(\mathbf{x}) = 0.5$ .

**Q11:** What does this decision boundary  $f(\mathbf{x}) = 0.5$  look like for the input space  $\mathbf{x}$ ? Reconstruct it and present it in a plot.

## Binary classification with logistic regression

We can turn our continuous regression problems into binary/categorical classification problems by mapping the outputs of our models to probabilities of a specific class. For the binary classification problem, we typically use a sigmoid function ( $\sigma(\cdot)$ ) that squeezes its inputs between 0 and 1; if the probability  $> 0.5$ , the outcome is class 1, and it is 0 otherwise.

**Q12:** Which of the nonlinearities you know would be useful for categorical (multi-class) classification problems and why?

For the sake of simplicity, we first consider a linear model followed by a sigmoid function to convert the outputs into probabilities:

$$p = \sigma(f(\mathbf{x}; \mathbf{w})) = \sigma(\mathbf{w}^T \mathbf{x}). \quad (1.5)$$

Learning the optimal parameters  $\mathbf{w}$  for such a classification problem is called *logistic regression*. As for adaptive filters, we will do this in a gradient-descent fashion. We will therefore now turn to calculating the required gradients.

Since our outcomes are binary, 0 or 1, the negative log-likelihood cost function is obtained from the Bernoulli distribution; i.e. the binary cross-entropy between the predicted probabilities and the true outcomes (labels/targets):

$$J = - \sum_{i=0}^{m-1} y^{(i)} \log(p^{(i)}) + (1 - y^{(i)}) \log(1 - p^{(i)}), \quad (1.6)$$

where  $m$  is the amount of data samples,  $y^{(i)}$  is the  $i^{\text{th}}$  label, and  $p^{(i)}$  is the predicted probability that the outcome is 1. In our case,  $p^{(i)} = \sigma(f(\mathbf{x}^{(i)}; \mathbf{w}))$ .

**Q13:** Derive the gradient of  $J$  with respect to  $p$ , i.e.  $\partial_p J$ .

**Q14:** Derive the gradient of  $p$  with respect to  $f$ , i.e.  $\partial_f p$  (see Q7).

**Q15:** Derive the gradient of  $f(\mathbf{x}; \mathbf{w})$  with respect to  $\mathbf{w}$ , i.e.  $\partial_{\mathbf{w}} f(\mathbf{x}; \mathbf{w})$ . To which update (linear) adaptive filter does this step relate?

**Q16:** Use the chain rule to derive the gradient of the cost function with respect to the weights, i.e.  $\partial_{\mathbf{w}} J$

## Classification with a shallow nonlinear model

We will now extend our logistic regression classifier with the more flexible nonlinear model we described earlier (Eqn. 1.4), i.e.

$$p^{(i)} = \sigma \left( f^{(2)} \left( f^{(1)}(\mathbf{x}; \mathbf{W}^{(1)}, \mathbf{b}^{(1)}); \mathbf{w}^{(2)}, b^{(2)} \right) \right) \quad (1.7)$$

**Q17:** Similar to Q13-Q16, use the chain rule to derive the gradient of the cost function with respect to all the model weights and biases, i.e.  $\partial_{\mathbf{w}^{(2)}} J$ ,  $\partial_{\mathbf{W}^{(1)}} J$ ,  $\partial_{b^{(2)}} J$  and  $\partial_{b^{(1)}} J$ .

This is a multiple choice question, with one possible answer for each of the four partial derivatives.

$$\begin{aligned} A : & \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} < 0 \\ (\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)}) & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ \partial_{\mathbf{w}^{(2)}} J = B : & \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} < 0 \\ -(\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)}) & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ C : & (\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)}) \\ D : & -(\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)}) \end{aligned} \quad (1.8)$$

$$\begin{aligned} A : & \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} < 0 \\ -(\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{w}^{(2)})^T \mathbf{x}^{(i)} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ \partial_{\mathbf{W}^{(1)}} J = B : & \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} < 0 \\ (\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{w}^{(2)})^T \mathbf{x}^{(i)} & \text{if } \mathbf{W}^{(1)} \mathbf{x}^{(i)} + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ C : & \sum_i \left( \frac{1-y^{(i)}}{1-\sigma(f^{(2)})} - \frac{y^{(i)}}{\sigma(f^{(2)})} \right) (\mathbf{w}^{(2)})^T \mathbf{x}^{(i)} \\ D : & -\sum_i \left( \frac{1-y^{(i)}}{1-\sigma(f^{(2)})} - \frac{y^{(i)}}{\sigma(f^{(2)})} \right) (\mathbf{w}^{(2)})^T \mathbf{x}^{(i)} \end{aligned} \quad (1.9)$$

$$\begin{aligned}
A &: 1 \\
\partial_{b^{(2)}} J &= \begin{aligned} B &: \begin{cases} \mathbf{0} & \text{if } \mathbf{x}^{(i)} < 0 \\ -\sum_i y^{(i)} - \sigma(f^{(2)}) & \text{if } \mathbf{x}^{(i)} \geq 0 \end{cases} \\ C &: -\sum_i y^{(i)} - \sigma(f^{(2)}) \\ D &: \sum_i y^{(i)} - \sigma(f^{(2)}) \end{aligned} \end{aligned} \tag{1.10}$$

$$\begin{aligned}
\partial_{\mathbf{b}^{(1)}} J &= \begin{aligned} A &: \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)}x + \mathbf{b}^{(1)} < 0 \\ (\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{w}^{(2)})^T & \text{if } \mathbf{W}^{(1)}x + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ B &: \begin{cases} \mathbf{0} & \text{if } \mathbf{W}^{(1)}x + \mathbf{b}^{(1)} < 0 \\ -(\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{w}^{(2)})^T & \text{if } \mathbf{W}^{(1)}x + \mathbf{b}^{(1)} \geq 0 \end{cases} \\ C &: 1 \\ D &: (\sum_i y^{(i)} - \sigma(f^{(2)})) (\mathbf{w}^{(2)})^T \end{aligned} \end{aligned} \tag{1.11}$$

The model you have just evaluated is actually a shallow, single-layer, *neural network*. The gradient calculation strategy you derived is called *back-propagation*. As for adaptive filters, we can now update our weights as:  $\mathbf{w}_{n+1} = \mathbf{w}_n + \mu \partial_{\mathbf{w}} J$ , with  $\mu$  being some learning rate.